
MINTS: Testing Strict Motif-Detector Claims in Nucleotide Transformers

Anonymous Author(s)

Affiliation

Address

email

Abstract

Genomic foundation models often make biological labels linearly accessible, but decodability alone does not identify an internal mechanism. We study whether attention heads in nucleotide transformers can be justified as biological motif detectors under a strict mechanistic evidence standard. MINTS applies tokenization-aware QK/OV circuit analysis, residual-stream probing with controls, matched attention enrichment, activation patching, and sparse feature search to DNABERT-2 and a tested Nucleotide Transformer backend. Before making mechanistic claims, we report raw-sequence task baselines: GC-only AUROC ranges from 0.6361 to 0.9088, and 3–6-mer TF-IDF AUROC ranges from 0.7956 to 0.9406. DNABERT-2 layer-11 residual readouts achieve AUROC values of 0.9137, 0.9383, 0.8954, and 0.8847 on promoter and splice-site tasks, showing strong label decodability rather than a circuit-level explanation. Controls sharpen this interpretation: random-label and position-only probes are near chance; splice-site probes remain about 0.25 AUROC above GC-only controls; promoter probes are close to strong GC/k-mer baselines and should be interpreted cautiously. The primary strict test is negative: across 51,249 GM12878 CTCF sequences, no DNABERT-2 head meets the registered QK-to-motif criterion or matched attention-enrichment criterion. Batch activation patching finds promoter-TATA over-restoration, while the full splice-donor rerun no longer produces threshold-crossing restoration. These task-specific interventions do not establish a CTCF motif-detector circuit. The results provide a calibrated evidence standard for nucleotide-transformer interpretability and show where representation-level success falls short of a strict motif-detector claim.

1 Introduction

Transformer models for nucleotide sequences now perform well on regulatory and splice-site prediction tasks [8, 4, 7]. However, much of the usual evidence for biological understanding remains predictive or post hoc: attribution maps, motif visualizations, or probes can identify useful correlations, but they do not by themselves explain how internal computations produce predictions [1, 12]. This distinction is central for mechanistic interpretability, where a claim about a component should survive direct tests of what it reads, where it attends, what it writes, and whether intervention on it changes the relevant behavior [5, 10, 6].

Nucleotide transformers are a useful setting for this question. Known motifs give natural ground-truth hypotheses: a motif can specify where a component would need to attend, what feature it would need to read, and which perturbation should matter. They also expose a common interpretability risk. Attention maps, motif visualizations, and high probe scores can suggest a biological mechanism even when the model is using GC content, genomic-coordinate artifacts, distributed features, or another correlated signal. MINTS is designed to make the required evidence explicit.

The motifs used here are intentionally canonical. CTCF is a DNA-binding factor with a well-studied binding motif and public ENCODE peak calls in GM12878 cells, making it a natural strict test case for motif-local attention. The TATA box is a short promoter-associated sequence whose mutation gives a simple promoter perturbation. Splice donor and acceptor dinucleotides mark exon-intron boundaries and give compact sequence edits for task-level causal tests. These examples are not meant to exhaust regulatory biology; they are selected because a correct circuit-level claim should specify nucleotide support, token support, and an intervention target.

This paper contributes a calibrated evidence standard and a completed case study. MINTS separates three claims that are often conflated: residual decodability, task-level causal restoration under activation patching, and strict circuit-level motif detection. The primary strict test is CTCF: under the registered QK-alignment and matched-enrichment criteria, neither tested backend yields a CTCF attention head that should be called a strict motif detector. The promoter and splice analyses play a supporting role by showing that task labels can be decodable and that interventions can produce task-specific causal signals without establishing the stricter CTCF circuit claim.

2 Related Work and Positioning

Mechanistic interpretability of transformers has developed standard tools for decomposing attention heads, locating causal pathways, and tracing computational graphs [5, 10, 2]. QK/OV decomposition is especially relevant here because it separates where a head attends from what it writes. Activation patching then gives an intervention-based test of whether the component matters for a behavior [6]. MINTS applies these established tools to genomic transformers with explicit motif-level criteria, using ground-truth motifs to test both positive circuit evidence and possible overinterpretation.

Genomic deep-learning interpretability has a different tradition. Models such as DeepBind, SpliceAI, DNABERT, GPN, and Nucleotide Transformer show that neural sequence models learn useful biological regularities [1, 7, 8, 3, 4]. Many analyses in this area focus on motif visualization, attribution, saliency, or downstream probing. These are valuable, but they usually do not establish an internal algorithm. A saliency map may highlight a motif, and a probe may decode a label, while the model’s actual computation may depend on a distributed feature, a correlated composition signal, or a different component entirely.

MINTS is positioned between these literatures. It does not claim to solve genomic interpretability in general, and it does not treat a high downstream or probe score as mechanistic evidence. Instead, it asks a narrower question: when, if ever, is it justified to call a specific attention head a biological motif detector? This choice deliberately favors falsifiable criteria over narrative completeness. Negative results are therefore useful evidence, because they identify where an appealing mechanistic story is not supported by the implemented tests.

3 Methods and Evidence Criteria

Biology primer. Table 1 defines the biological and tokenization terms used in the evidence stack.

Transformer circuits. Let $\tau(x) = (t_1, \dots, t_T)$ tokenize a nucleotide sequence into vocabulary items. The initial hidden state is

$$H_{j,:}^{(0)} = E_{\text{tok}}[t_j, :] + P_{j,:}. \quad (1)$$

For head h in layer l , with hidden state $H^{(l-1)}$, the attention pattern is

$$A_h^{(l)} = \text{softmax} \left(\frac{H^{(l-1)} W_h^Q (W_h^K)^T (H^{(l-1)})^T}{\sqrt{d_k}} \right), \quad (2)$$

and the head output is $A_h^{(l)} H^{(l-1)} W_h^V W_h^O$. Following transformer-circuit notation [5], the QK circuit $W_h^Q (W_h^K)^T$ controls attention scores, while the OV circuit $W_h^V W_h^O$ controls what is written to the residual stream. For nucleotide models, tokenization matters: a motif position may correspond to one nucleotide token, a fixed k -mer token, or a learned BPE span.

Table 1: Glossary for mechanistic-interpretability readers.

Term	Meaning in this paper
Motif	A recurring DNA pattern associated with a biological function, such as transcription-factor binding or splice recognition.
PWM	Position weight matrix: a per-position nucleotide model used to score how well a DNA window matches a motif.
JASPAR	A public database of curated transcription-factor binding motifs; MINTS uses the MA0139.1 CTCF matrix.
CTCF	CCCTC-binding factor, a DNA-binding protein involved in chromatin organization and regulatory insulation.
TATA box	A short A/T-rich promoter element used here for promoter perturbation and patching tests.
Promoter	A regulatory DNA region near a gene start site that can influence transcription initiation.
Splice	Boundary signals at the start and end of introns; donor sites usually contain GT and acceptor sites usually contain AG in genomic DNA.
donor/acceptor	
Nucleotide token	A model input unit corresponding to one or more DNA characters after tokenization.
k -mer	A fixed-length DNA substring of length k , such as a 6-mer token.
BPE	Byte-pair encoding, a learned variable-length tokenizer; DNABERT-2 BPE tokens can cover different numbers of nucleotides.

Table 2: Registered evidence criteria for candidate motif circuits. Thresholds are conservative screens: they are designed to prevent upgrading weak correlations, near-background attention, or unstable patching effects into motif-detector claims.

Claim	Statistic	Threshold	Purpose and scale
Motif localization	Attention enrichment ratio ρ_h	$\rho_h \geq 2.0$	Requires at least doubled motif-support attention over matched background.
Motif ranking	QK/motif Pearson correlation	$r \geq 0.5, p < 0.05$	Requires a large effect size, not merely significance from millions of tokens.
Causal importance	Patching metric $PM(h, l)$	$PM \geq 0.5$	Requires at least half restoration of the clean-minus-corrupted target effect.
Strong causal importance	Same	$PM \geq 0.75$	Marks stronger recovery while leaving room below exact full restoration.
Residual decodability	Probe AUROC or macro- F_1	≥ 0.80 or $+0.10$	Establishes diagnostic readout quality, not circuit use by the model.
Robustness	Matched negative controls	No matched artifact	Requires simple controls not to explain the candidate signal.

81 **Tokenization-aware motif support.** The strict CTCF tests use nucleotide coordinates as ground
82 truth and then map them to model-token coordinates. Let a motif hit span the half-open nucleotide
83 interval $[a, b)$ and let token j span $[u_j, v_j)$. Token j supports the motif when

$$\max(0, \min(v_j, b) - \max(u_j, a)) \geq 1, \quad (3)$$

84 that is, when the token overlaps the motif interval by at least one nucleotide. Special tokens
85 with zero-width offsets are retained for hidden-state alignment but receive no motif support. This
86 rule is implemented in `src/motif_scoring.py` and recorded in token-score manifests as
87 `token_support_min_overlap_bp=1`.

Sequence-to-token support mapping.

1. Score each nucleotide window with the JASPAR CTCF PWM and mark motif windows above the deterministic support threshold.
2. Convert each motif window to a half-open nucleotide interval $[a, b)$.
3. Tokenize the same sequence with offsets, retaining zero-width special tokens for alignment.
4. Mark every token whose nucleotide interval overlaps $[a, b)$ by at least one base.
5. Use the resulting token-support intervals for QK/motif correlation and matched attention-enrichment tests.

89 **Strict motif-detector criterion.** An attention head is treated as a candidate biological motif detector
90 only when three signals converge. It must assign elevated attention mass to motif-support positions,
91 its QK scores must rank motif-support tokens above matched background tokens, and patching its
92 activation must restore the task-relevant scalar in motif-destroying counterfactuals. The registered
93 thresholds used here are in Table 2. These thresholds intentionally make attention alone insufficient.

94 **Activation patching.** For a clean sequence x_c and motif-destroyed corrupted sequence x_{corr} , we
95 patch a head output from the clean run into the corrupted run. The restoration metric is

$$PM(h, l) = \frac{f(x_{\text{patched}}) - f(x_{\text{corr}})}{f(x_c) - f(x_{\text{corr}})}, \quad (4)$$

96 where f is a scalar target such as a probe decision function or class logit. $PM = 0$ indicates
97 no recovery, $PM = 1$ indicates full recovery of the clean effect, and values above one indicate
98 over-restoration rather than simple full recovery [6].

99 **Residual probes and controls.** Linear probes are trained on frozen mean-pooled residual vectors.
100 Their role is diagnostic: a high probe score establishes linear decodability, not that the probe direction
101 is causal or used by the model. We therefore compare residual probes with GC-content-only classifiers,
102 position-only classifiers where coordinates are available, GC-matched test subsets, shuffled-label
103 probes, and GC-content distribution shifts. This framing follows linear-representation work while
104 preserving the weaker interpretation that probes support [11].

105 4 Experimental Setup

106 The primary model is zhihan1996/DNABERT-2-117M, an encoder-style ge-
107 nomic transformer with learned BPE tokenization. The comparison backend is
108 InstaDeepAI/nucleotide-transformer-v2-100m-multi-species, used as a
109 tested fixed-6mer Nucleotide Transformer checkpoint [4]. The run exports DNABERT-2 QK and
110 OV matrices for all 12 layers and 12 heads per layer, caches residual vectors for layers 0, 5, and 11,
111 and uses custom Hugging Face forward hooks for attention-output patching because the installed
112 TransformerLens path does not expose all required DNABERT-2 per-head tensors.

113 The residual-probe tasks are TATA promoter, non-TATA promoter, splice donor, and splice acceptor
114 classification. The configured splits contain 5,062/212 train/test examples for TATA promoters,
115 30,000/1,372 for non-TATA promoters, and 30,000/3,000 for each splice task. The strict CTCF scan
116 uses 51,249 prepared GM12878 CTCF peak sequences and the JASPAR MA0139.1 CTCF motif,
117 aligned to exact overlapping DNABERT-2 BPE token spans. Batch patching uses motif-destroying
118 clean/corrupted pairs for promoter-TATA and splice-donor tasks.

119 To avoid treating probes as end-to-end model performance, we report raw-sequence predictive
120 baselines and a frozen DNABERT-2 sequence head before the mechanistic evidence. The raw-
121 sequence baselines are a GC-content logistic classifier using GC fraction, GC skew, and sequence
122 length, and a TF-IDF 3–6-mer logistic classifier. The frozen DNABERT-2 sequence head is a
123 balanced logistic classifier trained on cached layer-11 sequence embeddings. Frozen residual readouts
124 are shown beside these baselines only as representation context; neither DNABERT column is full
125 encoder fine-tuning.

126 The pipeline writes small review-facing CSV, JSON, and figure artifacts while treating large activation,
127 circuit-matrix, token-score, and sparse-autoencoder files as reproducible generated outputs. The June
128 2026 completed results required about 14.15 non-overlapping compute hours on the tested CUDA
129 setup after excluding a discarded overlapping cross-model attempt. The most expensive stages are
130 systematic causal intervention, cross-model comparison, and strict all-layer CTCF scans, so the
131 appendix includes both full reproduction commands and a reduced smoke-run command for checking
132 the code path on a smaller subset.

133 5 Results

134 **Predictive performance context before mechanism.** Table 3 separates task predictability from
135 mechanistic evidence. The raw-sequence baselines are strong, especially on promoter tasks: a
136 simple GC-content classifier already reaches 0.8955–0.9088 AUROC, and a 3–6-mer classifier
137 reaches 0.9297–0.9406 AUROC. Splice tasks are less explained by composition alone, but k-mer
138 baselines still recover substantial signal. The frozen DNABERT-2 layer-11 sequence head reaches
139 0.8847–0.9383 AUROC. The residual readout is included to show that the pretrained representation
140 is predictive, but it remains a residual-decoding result rather than a circuit explanation.

141 **Residual decodability is strong but not uniformly clean.** DNABERT-2 layer-11 residual probes
142 exceed the registered residual-decodability threshold for all four tasks. Table 4 combines headline
143 AUROC values with the most important controls. Random-label probes collapse to chance, and
144 position-only metadata does not explain the results. The splice-site tasks survive stronger checks:
145 donor and acceptor residual probes remain about 0.25 AUROC above GC-only baselines on matched

Table 3: Predictive task context before mechanistic analysis. GC-only and 3–6-mer rows are raw-sequence baselines. The frozen DNABERT-2 sequence head is trained on cached layer-11 sequence embeddings; the residual readout is reported separately as diagnostic decodability evidence.

Task	Train	Test	GC AUROC	GC AUPRC	GC Acc.	k-mer AUROC	k-mer AUPRC	k-mer Acc.	Seq-head AUROC	Readout AUROC
promoter_tata	5,062	212	0.8955	0.9040	0.8019	0.9297	0.9378	0.8255	0.9137	0.9137
promoter_no_tata	30,000	1,372	0.9088	0.9218	0.8316	0.9406	0.9482	0.8652	0.9383	0.9383
splice_sites_donors	30,000	3,000	0.6560	0.6892	0.6013	0.8185	0.8093	0.7430	0.8954	0.8954
splice_sites_acceptors	30,000	3,000	0.6361	0.6343	0.5883	0.7956	0.7792	0.7220	0.8847	0.8847

Table 4: DNABERT-2 layer-11 probe-control summary. Numeric columns report AUROC; random-label values are means over eight shuffled-label runs.

Task	Resid.	GC-only	Pos.-only	GC-match	Rand.	GC-shift range
promoter_tata	0.9137	0.8956	0.4136	0.9137	0.4700	0.6182–0.8030
promoter_no_tata	0.9383	0.9088	0.3865	0.9383	0.5129	0.6312–0.8368
splice_sites_donors	0.8954	0.6560	0.4414	0.8944	0.5064	0.8432–0.8669
splice_sites_acceptors	0.8847	0.6361	0.4461	0.8838	0.5005	0.8573–0.8580

subsets and remain strong under GC shifts. The promoter tasks require caution because GC-only baselines are already high; their residual-probe advantage over GC-only controls is small on the matched comparison.

This distinction changes the interpretation of the headline probe table. The splice-site result is evidence that DNABERT-2’s frozen residual stream contains linearly accessible information beyond simple GC composition or coordinate metadata. The promoter result is weaker mechanistically: it remains a valid decodability result, but the controls show that promoter-associated composition can explain much of the signal. This is why the paper avoids claiming that the promoter probes have isolated a clean causal promoter motif direction.

The strict CTCF motif-detector test fails. The CTCF scan produced 281,915 motif-support tokens and 1,843,874 finite token positions after BPE span correction. Across all 144 DNABERT-2 heads, the best QK-to-motif correlation is $r = 0.3004$, below the registered $r \geq 0.5$ criterion. The best matched attention-enrichment ratio is $\rho_h = 1.3130$, below the registered $\rho_h \geq 2.0$ criterion. The tested Nucleotide Transformer backend also fails to recover a strict CTCF candidate, with best $r = 0.0192$ and best $\rho_h = 1.00009$. Thus, the tested fixed-6mer backend does not support the CTCF motif-local head hypothesis under the same criteria.

The negative CTCF result is not a failure of the pipeline; it is the point of using a registered evidence stack. With more than 1.8 million finite token positions, small correlations can be statistically nonzero while remaining too small to support a mechanistic claim. The all-layer enrichment scan similarly prevents a common attention-overinterpretation mistake: a head can have visible structure without assigning enough mass to matched motif-support positions to count as a motif-local detector.

Threshold sensitivity supports the conservative conclusion. Table 6 and Figure 2 summarize threshold sweeps over QK correlation, matched enrichment, and auxiliary patching metrics. Relaxing the QK threshold alone yields 31 heads at $r \geq 0.1$, five at $r \geq 0.2$, one at $r \geq 0.3$, and none at $r \geq 0.4$. Relaxing matched enrichment alone yields three heads at $\rho_h \geq 1.1$, one at $\rho_h \geq 1.25$, and none at $\rho_h \geq 1.5$. The strict CTCF chain remains negative: no head jointly passes QK and enrichment once $r \geq 0.2$, even with $\rho_h \geq 1.1$, and the most permissive joint count at $r \geq 0.1$, $\rho_h \geq 1.1$ is two heads, equal to the 95th percentile of a permuted-head alignment null. Token-level motif-score nulls provide an additional calibration: the observed support-token count is 256,918, compared with a shuffled-score 95th percentile of 36,093.05 and zero nearest-GC non-support tokens above the motif-support threshold. The promoter and splice patching thresholds are more permissive because those are auxiliary task-specific interventions, not CTCF motif-detector proof.

Auxiliary promoter/splice patching gives task-specific causal signals. Batch promoter-TATA patching over 327 usable pairs identifies layer 4 head 8 with mean $PM = 1.4029$; layer 6 head 8 and layer 2 head 7 are also high. Because $PM > 1$ overshoots the clean-minus-corrupted effect, this is a methodologically sensitive causal signal rather than a simple full-recovery story. Batch splice-donor patching over 13,303 usable pairs identifies layer 7 head 0 as the largest restoration, but its mean

Table 5: Target-alignment ledger. CTCF is the primary strict motif-detector chain. Promoter and splice results are auxiliary evidence for task-specific residual decodability and causal signal, not substitutes for a CTCF proof.

Target	Performance	Probe	QK	Enrichment	Patching	OV	SAE	Role
CTCF	51,249-sequence strict scan; no classifier target	No CTCF residual classifier used as strict evidence	Best $r = 0.3004$; fail	Best $\rho_h = 1.3130$; fail	No candidate promoted after failed QK/enrichment	No aligned OV claim	Best cosine 0.1092; weak	Primary strict chain
Promoter-TATA	GC 0.8955; seq-mer 0.9297; head/readout 0.9137 AUROC	k- Decodable, but close to strong baselines	Not aligned current fact	target- Not aligned in current fact	Best PM = 1.4029; over-restores	L2H7 out- put cosine 0.1261; weak	Not promoter-specific	Auxiliary causal evidence
Splice donor	GC 0.6560; seq-mer 0.8185; head/readout 0.8954 AUROC	k- Decodable beyond GC and k-mer	be- Not aligned current fact	target- Not aligned in current fact	Best PM = 0.1657; below causal threshold	Not audited	Not splice-specific	Auxiliary negative-control evidence

Table 6: Threshold sensitivity summary. Counts are heads passing each single-metric threshold, except the joint row, which requires both CTCF QK and enrichment. Null references include permuted head alignment plus token-level motif-score and GC-matched background calibration.

Target	Sweep	Permissive count	Registered/strict count	Best value	Null reference
CTCF	QK r	31 at $r \geq 0.1$	0 at $r \geq 0.5$	0.3004	Motif-score shuffle: 256,918 observed vs 36,093.05 p95
CTCF	Enrichment ρ_h	3 at $\rho_h \geq 1.1$	0 at $\rho_h \geq 2.0$	1.3130	Position-matched background mean $\rho_h \approx 1.00$
CTCF	Joint QK/enrichment	2 at $r \geq 0.1, \rho_h \geq 1.1$	0 at $r \geq 0.5, \rho_h \geq 2.0$	-	Permuted-head 95th percentile is 2 at the permissive setting
Promoter-TATA	Patching PM	40 at PM ≥ 0.25	16 at PM ≥ 0.75	1.4029	Auxiliary task-specific signal
Splice donor	Patching PM	0 at PM ≥ 0.25	0 at PM ≥ 0.75	0.1657	Auxiliary task-specific negative result

PM = 0.1657 is below the causal threshold and the mean across heads is 0.0062. These interventions do not identify a strict CTCF motif detector because the CTCF QK and enrichment screens have no passing head.

OV, SAE, and cross-model audits remain conservative. The TATA OV readout audit for layer 2 head 7 finds only weak alignment with the trained TATA probe direction: the top output-write singular-vector cosine is 0.1261 and the probe self-gain through the OV matrix is -0.0326 . Sparse-autoencoder CTCF feature search is also weak across all 51,249 CTCF sequences, with best residual and MLP motif cosines of 0.1092 and 0.0906. In the protocol-level cross-model comparison, DNABERT-2 residual probes outperform the tested Nucleotide Transformer backend on the same tasks, but this is a checkpoint- and protocol-specific result, not a general claim about tokenization families or all Nucleotide Transformer variants.

6 What This Teaches Mechanistic Interpretability

The main lesson is that representation-level success and component-level mechanism are different claims. A residual probe can reveal that a biological label is linearly accessible while leaving open whether the model uses that direction, whether the feature is distributed, or whether the signal is driven by composition. Likewise, attention and motif visualizations can suggest a biological story without showing that a head’s QK circuit ranks motif-support tokens or that intervention on the head changes the relevant behavior.

MINTS treats biological motifs as useful test cases precisely because they provide external ground truth. A motif defines nucleotide support, gives a concrete perturbation target, and lets us ask whether model-token support agrees with circuit behavior. The negative CTCF result is therefore informative rather than merely null: it shows that strict evidence stacks can prevent a plausible but unsupported motif-detector story. For mechanistic interpretability, the practical recommendation is to report claims at the level they are actually supported: decodability, task-specific causal signal, motif-local attention, or strict circuit-level motif detection.

7 Limitations

This study is intentionally scoped to small and medium genomic transformers and does not claim exhaustive explanation. Current large-model circuit-tracing work also emphasizes partial recovery of

Table 7: Evidence ledger. Each row states the claim, the required evidence standard, the observed result, the supported interpretation, and the main limitation.

Claim	Required evidence	Observed result	Interpretation	Limitation
Promoter and splice labels are decodable from DNABERT-2 residual states	Strong residual readout plus controls against trivial artifacts	Layer-11 AUROC range 0.8847–0.9383; random-label and position controls near chance	Residual states contain task-relevant information	Promoter tasks remain close to strong GC and k-mer baselines
Splice residual decodability is not explained by GC alone	Matched and shifted GC controls should preserve the readout advantage	Splice probes remain about 0.25 AUROC above GC-only matched controls	Stronger evidence for non-composition residual signal on splice tasks	Still a probe result, not a localized circuit
Promoter-TATA patching identifies task-specific causal signal	Head-level patching should restore a corrupted scalar across many pairs	Best mean PM = 1.4029 over 327 usable pairs	A head can causally amplify the promoter-TATA readout under patching	Over-restoration and probe geometry prevent a simple full-recovery claim
Splice-donor patching does not identify a threshold-crossing causal head in the full rerun	Patching should restore a corrupted splice-donor scalar target	Best mean PM = 0.1657 over 13,303 usable pairs	The splice-donor intervention is weak despite strong residual decodability	Effect is small on average and not tied to QK/enrichment motif evidence
DNABERT-2 has a strict CTCF motif-detector head	One head should pass QK/motif ranking, motif-local enrichment, and causal evidence	Best QK $r = 0.3004$; best $\rho_h = 1.3130$; no passing candidate	The strict single-head CTCF detector claim is not supported	Negative result is scoped to selected model, motif scorer, data, and thresholds
Relaxed CTCF thresholds do not rescue the strict claim	Sensitivity sweep should reveal whether conclusions depend on the registered cutoffs	No joint QK/enrichment pass once $r \geq 0.2$, even at $\rho_h \geq 1.1$	The negative CTCF conclusion is not an artifact of one conservative cutoff	Per-head shuffled QK-score nulls require raw per-token QK vectors
SAE and OV audits do not reveal a simple alternate CTCF circuit	Output-write or sparse features should align strongly with the biological target	TATA OV cosine 0.1261; CTCF SAE best cosine 0.1092	Auxiliary audits remain consistent with a distributed or absent simple motif circuit	Audits are not exhaustive circuit discovery




Figure 1: Mechanistic evidence heatmaps. Left: DNABERT-2 CTCF QK-to-motif correlations. Middle: matched attention-enrichment ratios. Right: promoter-TATA batch patching. The CTCF heatmaps show a negative strict-proof result, while the TATA patching heatmap shows task-specific causal over-restoration.

computation rather than complete mechanistic accounts [2, 9]. The strict CTCF failure applies to the selected checkpoints, GM12878 CTCF intervals, JASPAR MA0139.1 motif scorer, BPE span-alignment rule, selected causal targets, and registered thresholds. It is not an impossibility theorem for all genomic transformers, all CTCF representations, all cell types, all motif scorers, or all circuit decompositions.

Linear probes remain limited evidence. They can reveal decodable information while exploiting correlated sequence composition, split artifacts, or metadata. The added GC, position, matched, random-label, and shift controls reduce this risk, especially for splice-site tasks, but promoter decodability remains entangled with GC composition. Activation patching is causal but incomplete: head-level restoration does not characterize the full circuit, and over-restoration can reflect denominator sensitivity, probe geometry, or out-of-distribution patched states. The threshold-sensitivity table reduces concern that the CTCF conclusion depends on one arbitrary cutoff, but it does not replace new data-generating null experiments. Finally, biological annotations are incomplete, and a single JASPAR PWM cannot represent all relevant regulatory mechanisms.

8 Conclusion

MINTS turns a common interpretability temptation into a falsifiable test: before calling an attention head a biological motif detector, require aligned QK scores, motif-local attention enrichment, and




Figure 2: Threshold sensitivity for QK alignment, matched attention enrichment, auxiliary patching, and joint CTCF QK/enrichment counts. Individual relaxed screens admit a small number of heads, but the strict CTCF chain remains negative under registered thresholds and under joint relaxation once $r \geq 0.2$.

causal restoration. On DNABERT-2, promoter and splice-site labels are strongly linearly decodable from late residual states, and splice-site decodability survives simple composition and position controls. But the strict CTCF motif-detector proof fails across all tested heads, and the task-specific patching effects do not change that conclusion. The contribution is therefore both positive and negative: a reproducible evidence standard for nucleotide-transformer interpretability, and a concrete demonstration that strong representation-level decodability should not be upgraded into a circuit-level biological claim without stronger evidence.

References

- [1] Babak Alipanahi, Andrew Delong, Matthew T. Weirauch, and Brendan J. Frey. Predicting the sequence specificities of DNA- and RNA-binding proteins by deep learning. *Nature Biotechnology*, 33(8):831–838, 2015. doi: 10.1038/nbt.3300.
- [2] Emmanuel Ameisen, Jack Lindsey, Adam Pearce, Wes Gurnee, Nicholas L. Turner, Brian Chen, Craig Citro, David Abrahams, Shan Carter, Basil Hosmer, Jonathan Marcus, Michael Sklar, Adly Templeton, Trenton Bricken, Callum McDougall, Hoagy Cunningham, Thomas Henighan, Adam Jermyn, Andy Jones, Andrew Persic, Zhenyi Qi, T. Ben Thompson, Sam Zimmerman, Kelley Rivoire, Thomas Conerly, Chris Olah, and Joshua Batson. Circuit tracing: Revealing computational graphs in language models. *Transformer Circuits Thread*, 2025. URL <https://transformer-circuits.pub/2025/attribution-graphs/methods.html>.
- [3] Gonzalo Benegas, Sanjit Singh Batra, and Yun S. Song. DNA language models are powerful predictors of genome-wide variant effects. *Proceedings of the National Academy of Sciences*, 120(44):e2311219120, 2023. doi: 10.1073/pnas.2311219120.
- [4] Hugo Dalla-Torre, Liam Gonzalez, Javier Mendoza-Revilla, Nicolas Lopez Carranza, Adam Henryk Grzywaczewski, Francesco Oteri, Christian Dallago, Evan Trop, Bernardo P.

- 251 de Almeida, Hassan Sirelkhatim, Guillaume Richard, Marcin Skwark, Karim Beguir, Marie
252 Lopez, and Thomas Pierrot. Nucleotide transformer: building and evaluating robust foun-
253 dation models for human genomics. *Nature Methods*, 22(2):287–297, 2025. doi: 10.1038/
254 s41592-024-02523-z.
- 255 [5] Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann,
256 Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep
257 Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt,
258 Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and
259 Chris Olah. A mathematical framework for transformer circuits. *Transformer Circuits Thread*,
260 2021. <https://transformer-circuits.pub/2021/framework/index.html>.
- 261 [6] Stefan Heimersheim and Neel Nanda. How to use and interpret activation patching. *arXiv*
262 *preprint arXiv:2404.15255*, 2024. doi: 10.48550/arXiv.2404.15255.
- 263 [7] Kishore Jaganathan, Sofia Kyriazopoulou Panagiotopoulou, Jeremy F. McRae, Siavash Fazel
264 Darbandi, David Knowles, Yang I. Li, Jack A. Kosmicki, Juan Arbelaez, Wenwu Cui, Grace B.
265 Schwartz, Eric D. Chow, Efstathios Kanterakis, Hong Gao, Amirali Kia, Serafim Batzoglou,
266 Stephan J. Sanders, and Kyle Kai-How Farh. Predicting splicing from primary sequence with
267 deep learning. *Cell*, 176(3):535–548.e24, 2019. doi: 10.1016/j.cell.2018.12.015.
- 268 [8] Yanrong Ji, Zhihan Zhou, Han Liu, and Ramana V. Davuluri. DNABERT: pre-trained bidi-
269 rectional encoder representations from transformers model for DNA-language in genome.
270 *Bioinformatics*, 37(15):2112–2120, 2021. doi: 10.1093/bioinformatics/btab083.
- 271 [9] Jack Lindsey, Wes Gurnee, Emmanuel Ameisen, Brian Chen, Adam Pearce, Nicholas L. Turner,
272 Craig Citro, David Abrahams, Shan Carter, Basil Hosmer, Jonathan Marcus, Michael Sklar,
273 Adly Templeton, Trenton Bricken, Callum McDougall, Hoagy Cunningham, Thomas Henighan,
274 Adam Jermy, Andy Jones, Andrew Persic, Zhenyi Qi, T. Ben Thompson, Sam Zimmerman, Kel-
275 ley Rivoire, Thomas Conerly, Chris Olah, and Joshua Batson. On the biology of a large language
276 model. *Transformer Circuits Thread*, 2025. URL [https://transformer-circuits.](https://transformer-circuits.pub/2025/attribution-graphs/biology.html)
277 [pub/2025/attribution-graphs/biology.html](https://transformer-circuits.pub/2025/attribution-graphs/biology.html).
- 278 [10] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing fac-
279 tual associations in GPT. In *Advances in Neural Information Processing Systems*, 2022.
280 arXiv:2202.05262.
- 281 [11] Kiho Park, Yo Joong Choe, and Victor Veitch. The linear representation hypothesis and the
282 geometry of large language models. In *Proceedings of the 41st International Conference on*
283 *Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 39643–
284 39666, 2024. URL <https://proceedings.mlr.press/v235/park24c.html>.
- 285 [12] Haopeng Yu, Heng Yang, Wenqing Sun, Zongyun Yan, Xiaofei Yang, Huakun Zhang, Yiliang
286 Ding, and Ke Li. An interpretable RNA foundation model for exploring functional RNA motifs
287 in plants. *Nature Machine Intelligence*, 6:1616–1625, 2024. doi: 10.1038/s42256-024-00946-z.

A Theoretical Details

Definition 1 (Biological Motif Detector). *An attention head h in layer l functions as a biological motif detector for motif m if there exists a nonempty set of relevant query positions $\mathcal{Q}_m$ and a threshold τ such that, for each $i \in \mathcal{Q}_m$,*

$$\mathbb{E} \left[\sum_{j \in \text{supp}_m(x)} A_h^{(l)}[i, j] \mid \text{supp}_m(x) \neq \emptyset \right] \geq \tau \quad (5)$$

and the corresponding expectation conditioned on $\text{supp}_m(x) = \emptyset$ is below τ . Here $\text{supp}_m(x)$ is the set of key positions whose local sequence windows instantiate motif m .

Theorem 1 (Layer-1 Motif Detection via QK Circuits). *Consider a layer-1 attention head operating on $H^{(0)}$. Let $s_m(j)$ denote a scalar motif score associated with key position j . If for a fixed query position i there are constants $a_i \in \mathbb{R}$ and $\gamma > 0$ such that*

$$q_i^T k_j = a_i + \gamma s_m(j) \quad (6)$$

for all admissible key positions j , then the induced attention distribution is monotonically increasing in $s_m(j)$. If motif-bearing positions are separated from background positions by a positive margin in s_m , the head is a motif detector for m .

Proof. The attention logit is $\ell_{ij} = q_i^T k_j / \sqrt{d_k} = a_i / \sqrt{d_k} + (\gamma / \sqrt{d_k}) s_m(j)$. Since $\gamma / \sqrt{d_k} > 0$, scaling preserves the ordering of motif scores. The exponential and softmax are strictly monotone in each logit when other logits are fixed, so positions with larger motif scores receive larger attention weights. A positive separation margin therefore induces an attention-mass threshold that separates motif-bearing from background positions. $\square$

Lemma 1 (Approximate Extension Under Linear Decodability). *Suppose that at layer $l - 1$ there exists a vector u_m such that*

$$u_m^T h_j^{(l-1)} = s_m(j) + \epsilon_j, \quad |\epsilon_j| \leq \delta. \quad (7)$$

If a head's QK score is affine in this decodable direction and the clean motif/background score margin is greater than 2δ , then the attention ordering between motif and background positions is preserved.

Proof. For motif-bearing j_m and background j_b , the score difference is proportional to $s_m(j_m) - s_m(j_b) + \epsilon_{j_m} - \epsilon_{j_b}$. The perturbation term lies in $[-2\delta, 2\delta]$. A clean separation greater than 2δ keeps the sign of the score difference fixed, so softmax preserves the motif/background ordering. $\square$

B Implementation and Reproducibility Details

The completed June 2026 result set was reconstructed from two log files. The first run covered `write_config` through `distributed_feature_search` from 13:37:04 to 23:26:04 local time on June 17, 2026. A first `cross_model_tokenization_comparison` attempt then ran until 02:27:59 on June 18 and reached 49,000 of 51,249 DNABERT-2 CTCF sequences, but that work was discarded because the cross-model stage was rerun from the start. The completed cross-model rerun ran from 08:51:36 to 13:11:27 on June 18, 2026, and its manifest reports 15,591.074 seconds. Removing the discarded overlap gives 50,930.636 seconds of completed pipeline time, or 14h 08m 50.636s. The wall-clock span from first logged step to final completed rerun was 23h 34m 23s. The inspected result tree contains 136 files totaling about 5.06 GiB, including CSV/JSON summaries, figures, NPZ archives, and sparse-autoencoder checkpoints.

The anonymized project repository is available at <https://anonymous.4open.science/r/MINTS/>. It contains the project code, configuration, local paper source, and reproduction instructions. The repository currently serves as the single source of truth for the submitted code; an explicit repository license has not yet been added.

The main reproduction commands are:

```

328 git clone https://anonymous.4open.science/r/MINTS/ MINTS
329 cd MINTS
330 python -m pip install -r requirements.txt
331 python main.py
332 python main.py --only-probe-controls

```

333 A smaller smoke run is:

```

334 python main.py --max-probe-train 512 --max-probe-test 256
335     --max-qk-alignment-sequences 128
336     --max-patching-pairs 32
337     --max-cross-model-qk-alignment-sequences 128
338     --max-feature-search-sequences 128 --sae-epochs 1

```

339 Important verification artifacts include:

```

340     • results/pipeline_run.json
341     • results/pipeline_run_cross_model_tokenization_comparison.
342       json
343     • results/tables/linear_probe_metrics.csv
344     • results/tables/linear_probe_controls.csv
345     • results/tables/downstream_task_performance.csv
346     • results/tables/target_alignment_table.csv
347     • results/tables/review_issue_matrix.csv
348     • results/tables/evidence_ledger.csv
349     • results/tables/threshold_sensitivity.csv
350     • results/qk_alignment/ctcf_qk_alignment.csv
351     • results/enrichment/ctcf_qk_alignment_matched_attention_
352       enrichment.csv
353     • results/patching/promoter_tata_batch_dnabert_activation_
354       patching.csv
355     • results/patching/splice_sites_donors_batch_dnabert_
356       activation_patching.csv

```

357 The heatmaps used in the paper are generated under results/figures/.

358 C Additional Results

359 The full configured-split residual probe AUROCs are 0.9137 for TATA promoters, 0.9383 for non-
360 TATA promoters, 0.8954 for splice donors, and 0.8847 for splice acceptors. The corresponding
361 AUPRC values are 0.9241, 0.9475, 0.9049, and 0.8954.

362 D Response-to-Review Checklist

363 Table 8 records the review issue matrix used to guide the revision. The anonymized reviewer IDs are
364 treated as issue sources rather than as claims about reviewer intent.

365 For promoter-TATA batch patching, 327 token-shape-preserving pairs were usable with no near-zero
366 denominator failures. The best mean score was layer 4 head 8 with PM = 1.4029, followed by layer
367 6 head 8 with PM = 1.3791 and layer 2 head 7 with PM = 1.3642. For splice donors, 13,303 pairs
368 were usable with no denominator failures, and the strongest mean restoration was layer 7 head 0 with
369 PM = 0.1657.

370 For the tested Nucleotide Transformer backend, residual-probe AUROCs were 0.6502, 0.6976,
371 0.5695, and 0.5647 on the same four tasks. DNABERT-2 therefore outperformed this backend by
372 0.2408–0.3259 AUROC in this protocol. This comparison should not be read as a general ranking of
373 tokenizers or genomic transformer families.

Table 8: Review issue matrix.

ID	Issue	Concern	Resolution	Status
mV6D	Clarity	Paper read as several analyses.	Abstract/introduction now state one calibrated evidence-standard claim; Results opens with task context and a target ledger.	Addressed
mV6D	Target mismatch	Promoter/splice evidence blurred into CTCF claims.	Promoter/splice patching is auxiliary; CTCF remains the strict chain.	Addressed
C5jL	Missing performance	Mechanistic claims lacked predictive context.	Added GC and 3–6-mer baselines, a frozen DNABERT sequence head, and frozen-readout context before mechanism.	Addressed
C5jL	MI contribution	Venue contribution needed sharpening.	Contribution is now a falsifiable standard for motif-detector claims.	Addressed
C5jL	Thresholds	Registered thresholds need scale/sensitivity.	Added threshold purposes, sweeps, and available null calibration.	Addressed
ndFG	Undefined biology	Motifs and biological terms need more support.	Added motif rationale, glossary, and nucleotide-to-token support rule.	Addressed
ndFG	Limited scope	Conclusions could overstate selected models/motifs.	Limitations and target ledger state the scoped negative CTCF result.	Addressed

Table 9: Response-to-review checklist with exact fixes.

Concern	Exact fix	Where it appears	Artifact
Writing clarity	Reframed the paper around one calibrated evidence-standard claim and one primary negative CTCF case study	Abstract, Introduction, target-alignment ledger, evidence ledger	results/tables/evidence_ledger.csv
Target mismatch	Separated CTCF strict motif-detector evidence from promoter/splice auxiliary decodability and patching evidence	Target-alignment table and evidence ledger	results/tables/target_alignment_table.csv
Missing performance context	Added GC and 3–6-mer raw-sequence baselines plus a frozen DNABERT sequence head before mechanistic claims, while keeping frozen readouts labeled as residual decodability	Predictive performance context subsection	results/tables/downstream_task_performance.csv
Weak MI contribution	Stated the contribution as a falsifiable standard for motif-detector claims rather than broad genomic interpretability	Introduction, Related Work, and “What This Teaches Mechanistic Interpretability”	Manuscript text
Threshold justification	Added threshold purposes, sensitivity sweeps, and available null calibration	Criteria table and threshold-sensitivity subsection	results/tables/threshold_sensitivity.csv
Undefined biology	Added motif rationale, glossary, and nucleotide-to-token overlap rule	Introduction and Methods glossary/mapping block	src/motif_scoring.py
Limited scope	Strengthened limitations to selected checkpoints, intervals, scorer, tokenization rule, causal targets, and thresholds; removed impossibility framing	Limitations and Conclusion	Manuscript text

E Declaration of LLM Usage

LLM-assisted tools were used during the project for mathematical exposition refinement, debugging support, and code-generation drafts through Codex with state-of-the-art GPT-5.4/GPT-5.5-class models. The research direction, core problem decomposition, step-by-step implementation decisions, and final scientific interpretation were directed by the author. The author reviewed the generated code, checked it for errors and failure modes, corrected mistakes, and takes responsibility for the final manuscript, code, and claims.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state the scope as an evidence stack for nucleotide-transformer interpretability, and the results distinguish residual decodability, causal patching, and the negative strict CTCF motif-detector test.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The paper includes a dedicated limitations section covering probe causality, GC confounding, patching over-restoration, checkpoint specificity, incomplete biological annotations, and compute limits.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The theoretical results state their assumptions in the theorem and lemma, with full proofs provided in Appendix A.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.

- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper describes the models, datasets, splits, controls, metrics, and result artifacts needed to reproduce the main empirical claims, with commands in Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The code and reproduction instructions are provided through the anonymized repository <https://anonymous.4open.science/r/MINTS/>; the appendix lists the main commands and verification artifacts.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The experimental setup reports model checkpoints, task splits, motif source, scan size, controls, and intervention settings needed to interpret the results.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Probe results include bootstrap confidence intervals in the reproduced artifacts, QK alignment reports statistical tests, and the paper states where effects are interpreted by magnitude rather than significance alone.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.

- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix B reports the full-run runtime and the computationally dominant pipeline stages.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: The author has reviewed the NeurIPS Code of Ethics. The work is non-clinical interpretability research using public genomic-model resources, does not involve human subjects, does not release a deployable biological or medical decision system, and transparently reports limitations, code availability, and LLM assistance.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper discusses possible implications for genomic-model interpretability and explicitly avoids clinical deployment claims.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.

- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The work does not release a deployable clinical system, biological intervention protocol, or model intended for direct decision-making.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No]

Justification: The paper credits the main existing models, datasets, and motif resources, but the manuscript and repository do not yet explicitly list every asset license and term of use.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [No]

Justification: The project repository is a new released asset and contains documentation and reproduction commands, but it currently does not include an explicit repository license, so the asset documentation is incomplete.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The work does not involve crowdsourcing or human-subject data collection.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The work does not involve human subjects research requiring IRB approval.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: The appendix declares that LLM-assisted coding and writing tools were used to refine mathematical exposition and generate implementation drafts, while the author originated the research direction, directed the problem solving, checked the code, corrected errors, and takes responsibility for the final work.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.